

VIJAY BHATT

Principal Architect | Distributed Systems | AI, RAG & Agentic Systems | Cloud-Native & Event-Driven Architecture

Pune, India | +91-9764834026 | vjbhatt760@gmail.com

AI Career Chatbot: github.com/Vijay2351989/ai-career-conversations — Ask an AI agent about my career instead of reading this resume!

PROFESSIONAL SUMMARY

Results-driven Principal Architect with 15+ years of experience designing and delivering production-grade distributed systems, cloud-native microservices, and AI-powered enterprise solutions. Deep expertise in Node.js, TypeScript, and Python, with a strong track record of architecting scalable event-driven systems and microservice ecosystems. Extensive experience implementing microservices design patterns including API Gateway, CQRS, Saga, circuit breaker, and service mesh to build resilient, loosely coupled distributed architectures. Skilled in building RAG pipelines, AI-powered Q&A systems, agentic workflows, and document understanding platforms, with hands-on experience designing agentic patterns such as ReAct agents, evaluator-feedback loops, parallel sub-agent orchestration, and hierarchical multi-agent systems. Proven ability to lead cross-functional engineering teams, mentor developers, and drive measurable impact in system design, operational efficiency, and platform modernization across domains including energy, logistics, collaboration, and procurement.

CORE SKILLS / TECHNICAL SKILLS

Programming Languages & Frameworks: JavaScript, TypeScript, Python, Node.js, Express, Angular, React, Redux, FastAPI

Backend: Node.js/Express, FastAPI, REST APIs

AI — Agentic: LangChain, LangGraph, LLM Integration, AI Agents, Memory Management, Agentic Design Patterns (ReAct, Evaluator-Feedback, Parallel Sub-Agent, Hierarchical, Reflection/Refinement, Stochastic Consensus Multi-Agent, Agent Chatroom), Observability/Evaluation (Langfuse, MLflow, LangSmith)

AI — LLM Testing, Evaluation & Tracing: DeepEval, RagaAI Catalyst, Hugging Face Evaluate, LLM-as-a-Judge Evaluation, Regression Testing for LLM/RAG Apps, Trace-Based Evaluation for Agentic and Multi-Agent Systems

AI — Open-Source Models & Inference: Hugging Face Transformers, Open-Source LLM Integration, Model Benchmarking/Evaluation Pipelines, vLLM Inference Serving, FastAPI Model APIs

AI — RAG Pipelines: AI Q&A Systems, Multi-Hop RAG, HyDE RAG, Hybrid Retrieval (BM25 + Semantic Search, Reciprocal Rank Fusion), Reranking (Bi-Encoder, Cross-Encoder), ANN Indexing (HNSW, FAISS), Embeddings, Vector Databases (Qdrant, ChromaDB, Pinecone), Document Understanding

Event Streaming & Messaging: Apache Kafka, Apache Artemis, Message Bus Architecture, Queues, Topics, Event-Driven Patterns

Databases: MongoDB, Redis, PostgreSQL

Architecture: Microservices, Event-Driven Architecture, Distributed Systems, API Gateway, CQRS, Saga, Circuit Breaker, Service Mesh, Micro Frontends (Module Federation)

DevOps & Containerization: Docker, Kubernetes, Helm, Jenkins, Travis CI, GitHub Actions, CI/CD Pipelines, Nginx, Blue/Green & Canary Deployments

Frontend: Angular, React, HTML5, CSS3, Webpack, Babel, Module Federation

Prompt Engineering: DSPy (Signatures, Modules for programmatic prompt construction), Structured Prompt Design (role definition, goal specification, constraints/rules, failure conditions), Few-Shot Prompting, Information-Dense Prompt Composition

Context Engineering: Context Summarization, Tiered Condensation, Memory Extraction, Long-Term Memory Management

AI Coding Tools: Claude Code (agents, reusable skills), Augment Code

PROFESSIONAL EXPERIENCE

Principal Architect | Krista AI | Aug 2024 – Present

Krista is an agentic platform that transforms businesses' operations by orchestrating people, systems, and AI using everyday language. Designed for simplicity and flexibility, Krista enables organizations to automate processes across multiple systems and streamline workflows with conversational ease—no coding required.

Agentic Workflow Engine — ReAct Pattern-Based Autonomous Agent

- Designed and architected an **agentic workflow engine** using the **ReAct (Reason-Act-Observe) pattern**, enabling users to provide natural language workflow prompts and select tools, with the agent autonomously iterating through decision, action, and observation cycles to complete complex workflows.
- Built the agent using **LangChain4j** with multi-LLM support (**OpenAI, Claude**), designing structured **tool definition schemas** for seamless tool invocation and response handling.
- Used **DeepEval** to test multi-agent tool-calling precision, creating **golden datasets** and **LLM-as-a-Judge** evaluation criteria for multi-turn agent testing, simulation, and regression evaluation; published and reviewed evaluation results in **Confident AI**.
- Implemented a **hierarchical agent architecture** with child agents for large-scope tools (e.g., extension invoker), preventing context window overpopulation in the main agent and improving accuracy and cost efficiency.
- Designed and implemented **guardrails** for agent behavior control, along with comprehensive **token usage tracking** and a **cost monitoring dashboard** for full observability of LLM consumption across workflows.
- Architected the **Policy Enforcer** system—a **multi-agent evaluation pipeline** where authors upload policy documents to a document set, and a dedicated **Policy Agent (ReAct)** assists in selecting relevant policy chunks during authoring. These policy constraints are enforced at runtime via an **evaluator-feedback child agent** that evaluates the main agent's tool selections against policy rules, approving or rejecting actions, forcing the main agent to select alternate tools or gracefully exit when policy constraints cannot be met.

Conversation Builder — Multi-Agent Human-in-the-Loop Workflow Authoring

- Architected a **chat-based workflow authoring experience** for Krista's agentic workflow platform, allowing users to describe workflows in natural language instead of manually stitching steps, step relationships, branching conditions, variables, and attributes through a plug-and-play UI.
- Designed a **phase-wise multi-agent workflow** using **LangGraph**, covering requirement gathering, step identification, workflow creation, workflow editing, and deterministic refinement, with **human-in-the-loop feedback** at each phase to correct and improve agent output.
- Built real-time **WebSocket** communication and resumable **session management**, persisting each builder session through **EFS-backed persistent volumes** so users could resume from any workflow phase.
- Set up a **Playwright-based performance testing framework** for concurrent sessions, and designed runtime scaling with **sticky sessions** plus **KEDA-driven HPA autoscaling** based on active-session load per pod.

AI Q&A Module — Retrieval Pipeline Optimization for Enterprise Document Sets

- Drove end-to-end optimization of the **AI Q&A module** used by clients to perform question-answering over large document sets (collections of enterprise documents ingested into a RAG pipeline), significantly improving retrieval accuracy and response relevance.
- Redesigned the retrieval pipeline starting with **query decomposition** using **HyDE (Hypothetical Document Embeddings)** combined with **multi-hop reasoning** to break complex queries into intermediate reasoning steps and generate richer search context, and introduced a **hybrid search strategy** combining **lexical (BM25)** and **semantic (vector) search** to capture both keyword-

exact and contextually similar results.

- Implemented **Reciprocal Rank Fusion (RRF)** to merge results from lexical and semantic retrieval into a unified ranking, followed by **Cross-Encoder reranking** to score and surface the most relevant chunks—delivering a measurable improvement in answer quality over the original single-strategy RAG approach.
- Optimized retrieval performance for growing document sets by introducing **metadata filtering** and **database indexing strategies**, enabling faster and more targeted chunk retrieval as document volume scaled.
- Currently evaluating **Approximate Nearest Neighbor (ANN)** techniques for sub-linear retrieval at scale—completed a **proof-of-concept on Qdrant vector database using the HNSW algorithm**, and evaluating **FAISS** for in-memory ANN indexing. Additionally exploring **knowledge graph-based storage and retrieval** for improved entity-relationship search across document sets (under development, not yet production-deployed).

Meeting Agent — Enterprise Conversational Memory Platform

- Served as **design architect** for the Meeting Agent project, transforming meeting interactions into a centralized, searchable enterprise memory—replacing disconnected recording silos with a unified cognitive resource.
- Architected and built a **meeting bot** that joins live meetings, scrapes underlying meeting platforms, and captures real-time transcripts via captions using **Playwright**, feeding data into the system through an **event-driven architecture**.
- Designed and implemented the **Conversation Agent**, an event-driven pipeline using **Apache Kafka** consumers to process meeting events, with optimized caching layers, feeding meeting data into a **RAG document store** for enterprise-wide Q&A and knowledge retrieval.
- Built a **meeting management system** with **OAuth-based authentication**, enabling users to view upcoming meetings, configure bot permissions, define **custom summary templates**, and control bot lifecycle—with the Conversation Agent autonomously managing bot orchestration and generating tailored meeting summaries.

Cross-Service Interoperability Framework — Polyglot Service Contract Architecture

- Designed and implemented an **enterprise-grade interoperability framework** enabling seamless communication between services written in different languages (Node.js, Python, Java) using **Protocol Buffers (Proto)** and **Apache Artemis** messaging.
- Architected a **common service contract project** where service definitions, including method signatures, input/output schemas, and annotations, were defined in a language-agnostic manner, with Proto schema generation for input/output POJOs.
- Built a **code generator layer** that transforms service definitions and Proto schemas into **language-specific client/server libraries** for Node.js, Python, and Java, enabling type-safe cross-language service communication.
- The generated libraries internally manage **Artemis connection handling**, **Proto schema registration**, and message serialization—allowing application teams to simply include the library as a dependency and use structured service definitions to communicate across the polyglot microservice ecosystem.

Python Infrastructure Library — Common Infra Components for Python Services

- Designed and built a **shared Python infrastructure library** as the organization transitioned from Java to Python for better AI/ML library support, providing standardized, best-practice wrappers for core infrastructure components to be adopted across all Python services.
- Implemented an **Apache Artemis wrapper** with special attention to **consumer architecture**—supporting both **client-acknowledge** and **auto-acknowledge** modes, **consumer backpressure handling**, configurable **consumer thread pools** for high-load scenarios, and key configurations including **queue durability**, **message TTL**, and delivery guarantees.
- Built an **Infinispan wrapper** with support for **Proto schema registration** enabling structured data storage in Protobuf format, configurable **message TTL** and **eviction strategies** (LRU, LIRS,

size-based) for cache lifecycle management.

- Implemented a **rate limiter module** supporting key algorithms including **Token Bucket**, **Sliding Window Log**, **Sliding Window Counter**, and **Fixed Window Counter**. Supported **key-based rate limiting** backed by **Infinispan** for distributed state, enabling **global rate limiting**, **per-user rate limiting**, and **concurrency control via semaphore**—all exposed as **Python decorators** for seamless integration into service endpoints.

Tech Lead | SLB (Schlumberger) | May 2023 – Aug 2024

PTS (Petro Technical Suite) is a key petrotechnical program in SLB responsible for provisioning cloud VMs for customers, with auto-deployment of applications on user login, abstracting cloud complexities and delivering ready-to-use environments for petrotechnical operations.

- Joined as **Technical Lead** in the shared services team, managing a suite of **Node.js microservices** including the core provisioning service that triggers VM-related operations on cloud based on user requests.
- Led **architectural improvement initiatives** across all shared services, handling enhancements, bug fixes, and modernization of legacy components.
- Designed and implemented a **billing microservice** from the ground up, enabling usage-based billing of customers for cloud VM consumption—a key business initiative for the program.
- Scaled billing APIs using **rate limiting and semaphore-based concurrency control**, and separated read/write data stores using the **CQRS pattern with event sourcing** for high-throughput billing operations.
- Implemented **Saga orchestration** pattern for managing distributed transactions across microservices, ensuring data consistency in complex multi-step provisioning and billing workflows.
- Introduced **circuit breaker patterns** for inter-service communication, improving fault tolerance and reliability across the microservice ecosystem.
- **Re-architected monolithic services** into lean, decoupled microservices with **polyglot persistence** (introducing purpose-fit databases per service) and deployed them on the internal **LightOps Kubernetes cluster** following latest organizational standards.
- Implemented **Azure Event Grid** integration for real-time capture of VM deployment status events, improving visibility and reliability of provisioning workflows.
- Led **sprint planning and requirement grooming** with the team, presenting technical proposals to senior stakeholders for requirement review and season planning.
- Mentored team members on implementation, ensuring timely delivery of key work items aligned with season goals.

Technology Lead | Bigstep Technologies | Oct 2021 – Apr 2022

- Served as **Technology Lead** working as a **full-stack MEAN developer**, responsible for designing and developing features across the stack using **MongoDB**, **Express**, **Angular**, and **Node.js**.
- Led technical decision-making, mentored team members, and contributed to architecture and implementation of application features.

Lead Engineer | ElasticRun | Apr 2022 – Apr 2023

Matrix is the in-house ERPNext platform for ElasticRun, supporting all HR and Accounts related functionalities at the organization level.

- Analyzed multiple applications under the **Matrix umbrella** and led the initiative to extract and consolidate core shared functionality into **Matrix Core**—a reusable framework adopted across teams.
- Coordinated with cross-functional teams to drive **reusability, standardization, and adoption** of the Matrix Core framework across multiple Matrix applications.
- Managed **deployment and optimization** of the Matrix Core framework for production environments.

- Redesigned and improved the complete **CI/CD strategy and deployment automation** pipeline for the application suite.

Senior Engineer | Barco | Mar 2020 – Sep 2021

Chimay is a cloud-serviced product enabling multisite collaborative meetings, supporting various communication platforms with a rich collaborative experience including whiteboarding and active content sharing via a message bus architecture.

- Led the **greenfield setup** of the application, establishing the **microservices architecture** foundation including service decomposition, inter-service communication patterns, and deployment strategy.
- Designed **RESTful API contracts** and **database schemas** across services, ensuring scalable and maintainable data models.
- Built backend **Node.js microservices** for core meeting features—meeting service, adapter services, and room pairing service.
- Managed **cloud deployments** including **Helm chart creation** for individual services and deployment to **Kubernetes clusters**.
- Implemented **zero-downtime deployments** using rolling update and **blue/green deployment strategies**.
- Implemented **application health monitoring** using NewRelic ping monitors and **application analytics** using StatsD, Graphite, and Grafana dashboards.
- Designed **message bus architecture** using PubNub for cascading notifications to frontend clients, including secure subscription and client presence detection.

Officer | Citicorp Services India Pvt Ltd | Dec 2018 – Mar 2020

Governance Risk Compliance is an Operational Risk module in banking, enabling bankers to forecast risk for operational activities across regions, apply control measures, and generate severity and mitigation metrics.

- Led the **micro frontend architecture** initiative using **Webpack Module Federation**, building container/shell applications hosting multiple independent micro frontend apps in Angular.
- Created **reusable UI component libraries** using Angular Elements and ngx-build, enabling cross-project component sharing across the organization.
- Led offshore team coordination with onshore stakeholders, driving **requirements gathering and design implementation**.
- Designed **Angular components, routes, and modules**, integrating third-party libraries like PrimeNG with custom styling to meet application requirements.
- Supported **SIT and UAT** cycles with timely defect resolution.

Product Engineer → Member of Technical Staff | Infosys / Edgeverve System Ltd | Dec 2011 – Dec 2018

Started career as a fresher and progressively grew from junior developer to senior developer and finally into a technical lead role over 7 years across two major product lines—ProcureEdge (B2B Procurement) and AssistEdge (RPA Automation).

ProcureEdge Cloud Platform (Dec 2011 – Jul 2017)

- Began career as a **fresher**, building foundational skills in **JavaScript, Node.js, HTML, CSS, and MongoDB** while working on an enterprise-grade B2B procurement platform.
- Gradually took on increased responsibility, progressing from implementation tasks to **requirement analysis, technical feasibility assessment, and effort estimation**.
- Grew into a **senior developer and mentor** role, leading junior developers with code reviews, deliverable tracking, and technical guidance.

- Worked on the **operational procurement module** including shopping cart, team cart, multi-step approval workflows, and supplier compliance features.
- Gained deep understanding of **end-to-end development lifecycle**—from FSD analysis and wire-frame interpretation to development, unit testing, and deployment.

AssistEdge Cloud Platform — RPA Automation (Aug 2017 – Dec 2018)

- Transitioned into a **technical staff member** role on Edgeverve's **Robotic Process Automation** product, competing with Automation Anywhere and UiPath.
- Worked on **Java enterprise application automation** (Oracle Fusion, Oracle EBS), hooking into Java processes using agents to identify and automate Swing/AWT UI components.
- Tackled complex technical challenges around **JDK compatibility, application security**, and cross-environment automation feasibility.

KEY PROJECTS

Meeting Agent — Enterprise Conversational Memory Platform (Krista AI)

Designed and architected an enterprise meeting intelligence platform that transforms video call interactions into a centralized, searchable organizational memory. Built a meeting bot using Playwright that joins live meetings, captures real-time transcripts, and feeds data through an event-driven architecture using Apache Kafka. Designed the Conversation Agent—an event-driven pipeline with Kafka consumers, optimized caching, and a RAG document store enabling enterprise-wide Q&A over meeting content. Built a meeting management system with OAuth authentication, configurable summary templates, bot lifecycle management, and meeting owner controls.

Agentic Workflow Engine — ReAct Pattern Autonomous Agent (Krista AI)

Architected an agentic workflow engine using the ReAct (Reason-Act-Observe) pattern with LangChain4j and multi-LLM support (OpenAI, Claude). Users provide natural language prompts and select tools; the agent autonomously iterates through decision-action-observation cycles to complete workflows. Designed hierarchical agent architecture with child agents for complex tools, preventing context overpopulation. Used DeepEval with golden datasets, LLM-as-a-Judge criteria, multi-turn simulation, and Confident AI result tracking to evaluate multi-agent tool-calling precision and regression quality. Implemented a Policy Enforcer system—a multi-agent evaluation pipeline where policy documents feed into a dedicated Policy Agent that enforces constraints on the main agent's tool selections via an evaluator-feedback child agent, approving or rejecting actions in real time.

Conversation Builder — Multi-Agent HITL Workflow Authoring (Krista AI)

Designed a chat-first workflow authoring system for Krista's agentic workflow platform, enabling users to describe workflows in natural language instead of manually stitching steps, branching conditions, step relationships, variables, and attributes through the UI. Built a phase-wise multi-agent flow using LangGraph across requirement gathering, step identification, workflow creation, workflow editing, and deterministic refinement, with human-in-the-loop feedback at each phase. Implemented WebSocket communication, EFS-backed persistent session storage for resumable workflow creation, sticky session handling, Playwright-based concurrent-session performance testing, and KEDA-driven HPA autoscaling based on active-session load per pod.

Common Agentic Framework — Enterprise Agent-as-a-Service Platform (Krista AI)

Identified a critical organizational gap where multiple teams were independently building AI agents with inconsistent design patterns, no standardized context or memory management, and unaddressed agent security concerns. Designed and built a centralized agentic framework that allows applications to define agent configurations—model selection with fallback, guardrail policies, token limits, and context management strategies—and register reusable prompt templates with static and dynamic segments (runtime injection of tool definitions, structured output schemas, and domain-specific data). Applications create sessions against the framework, selecting their configured agent type and template; the prompt builder assembles the final prompt, and the agent executes within a managed lifecycle. The framework handles conversation memory (cache-backed short-term buffer, database-backed long-term persistence), enforces

guardrails, and manages context windows per the application's selected strategy. Each application's system prompt is wrapped inside a framework-level agent system prompt that provides self-annealing behavior and graceful recovery semantics. Built with Python, LangGraph/LangChain, with cache for short-term and database for long-term memory management.

Open-Source Model Evaluation & GPU Inference Platform

Built hands-on expertise with open-weight model evaluation and serving using Hugging Face Transformers and Hugging Face Evaluate. Ran model evaluation pipelines across diverse datasets, measuring precision, recall, F1 score, and accuracy, and tracked experiment runs, metrics, and comparative results in MLflow. Hosted GPU-backed inference for Qwen, DeepSeek, Llama, and Gemma-family open-weight models using vLLM and FastAPI, and integrated Gemini-family APIs where appropriate, exposing reusable model APIs for text generation, reasoning, extraction, classification, and multimodal application workflows.

AI Q&A Module — Retrieval Pipeline Optimization (Krista AI)

Drove multi-layer optimization of the enterprise AI Q&A system used by clients for question-answering over large document sets. Replaced traditional single-strategy RAG with a hybrid retrieval pipeline: HyDE with multi-hop reasoning for richer query decomposition, combined BM25 lexical and semantic vector search with Reciprocal Rank Fusion for unified ranking, and Cross-Encoder reranking for top-k chunk selection—significantly improving answer quality. Optimized retrieval performance at scale using metadata filtering and database indexing. Currently evaluating ANN techniques (completed POC on Qdrant with HNSW, evaluating FAISS) and knowledge graph-based retrieval for entity-relationship search.

Cross-Service Interoperability Framework (Krista AI)

Designed a polyglot service contract framework enabling type-safe communication between Node.js, Python, and Java services using Protocol Buffers and Apache Artemis. Built a code generator that transforms language-agnostic service definitions and Proto schemas into language-specific client/server libraries with built-in Artemis connection handling and schema registration, adopted as a shared library across the microservice ecosystem.

Cloud VM Billing & Microservices Modernization (SLB)

Designed and implemented a usage-based billing microservice for cloud VM consumption within the PTS program. Applied CQRS with event sourcing to separate read/write data stores, scaled APIs using rate limiting and semaphore-based concurrency control, and implemented Saga orchestration for distributed transactions across provisioning and billing workflows. Led the re-architecture of monolithic shared services into lean, decoupled microservices with circuit breaker patterns, deployed on the internal LightOps Kubernetes cluster.

Matrix Core Framework (ElasticRun)

Led the extraction and consolidation of shared functionality from multiple ERPNext applications into a reusable Matrix Core framework. Drove cross-team adoption, standardization, and redesigned the complete CI/CD strategy and deployment automation pipeline for the application suite.

ProcureEdge Cloud Platform (Infosys/Edgeverve)

Worked on the operational procurement module of ProcureEdge—a cloud-based B2B Source-to-Pay platform for global enterprises. Implemented shopping cart, team cart, and multi-step configurable approval workflows, enabling enterprises to streamline procurement operations, reduce TCO, and improve supplier compliance.

PERSONAL PROJECTS

AI Career Conversations | [GitHub](#)

An AI-powered career chatbot that serves as an interactive virtual resume—enabling interviewers and recruiters to learn about my background, skills, and experience through natural conversation instead

of reading a static document. Built with OpenAI GPT-4o Mini and Gradio, the agent is contextually augmented with my resume and career summary, and includes custom tools for handling edge-case queries and push notifications for unanswered questions. Deployed on HuggingFace Spaces. *Beyond its utility, this project itself demonstrates my ability to design and deploy AI agents with domain-specific knowledge grounding.*

Claude Skills & Agents | [GitHub](#)

A curated collection of 14+ reusable Claude Code skills and agents for automating software development workflows and everyday tasks. Includes programming-focused agents (spec creation, systematic debugging, code review, project management, test architecture) and task automation agents (email management, movie recommendations, shopping). Built with Python, TypeScript, and Shell.

AI Agentic Memory Management | [GitHub](#)

A dual-layer memory system for LLM-powered chatbots that solves unbounded token growth in long conversations. Implements bounded short-term context using Redis with automatic summarization every N turns, and persistent long-term memory using Qdrant vector database for semantically searchable fact storage. A DSPy-based ReAct agent intelligently decides whether to add, update, or delete memories, while response generation and memory management operate as decoupled pipelines. Built with Python, DSPy, Qdrant, Redis, and OpenAI embeddings.

AI Worker Sidekick Agent | [GitHub](#)

An autonomous agent framework built with LangGraph implementing the worker-evaluator-optimizer design pattern. Features a Gradio chatbot interface with per-user session management, human feedback loops, and iterative optimization cycles. The worker agent executes tasks using integrated tools (web search, Wikipedia, file management), the evaluator validates output against success criteria, and the optimizer iteratively improves results. Deployed on HuggingFace Spaces.

AI Corrective RAG Agent | [GitHub](#)

A multi-agent corrective RAG system built with LangGraph that intelligently routes queries between a ChromaDB vector store and web search based on semantic similarity scores. Features a 5-stage pipeline: Planner (query refinement), Classifier (metadata-based categorization), Score (similarity threshold evaluation), RAG (vector retrieval), and Web Search (fallback). Uses semantic chunking and configurable thresholds.

AI Deep Research | [GitHub](#)

An autonomous deep research system built with the OpenAI Agent SDK featuring a multi-agent pipeline: Planner (generates search strategy with relevant terms), Search (executes web searches), Writer (synthesizes findings into comprehensive markdown reports), and Research Manager (orchestrates the workflow). Users provide a topic, confirm search terms, and receive structured research summaries. Deployed on HuggingFace Spaces.

EDUCATION

Bachelor of Technology | Institute of Engineering and Management | 2011 | DGPA: 8.57 / 10
West Bengal University of Technology (W.B.U.T.)